

SERVICES

About Services:- A grouping of pod endpoints into a single endpoint is called as services. (pod endpoints are defined as pod ipaddress) it will be used as load balancer. pod can be down at any time and we cannot assure the pod will come back with same ip address and it will be a dynamic assign, so to connect pod consistently, the services will create static endpoints. so, in backend ephemeral pods and endpoints will be combined and give as static endpoint by services, so static point will remain a constant until we remove, however backend whatever pods present that can be deleted or down at any time, that will be taken care by services. With Selectors and labels, pods will be grouping by services. we never bother about backend pods, we will be connecting with services, it will take care of all pods in backend. There are five types of Services (ClusterIP (default), NodePort, LoadBalancer, ExternalName, Headless). Ingress is top of services, so in real-time ingress will be used to expose outside.

Actual readme:-

- What is a Service?

The idea of a Service is to group a set of Pod endpoints into a single resource

- Why to use a Service?

In a Kubernetes cluster, each Pod has an internal IP address. But the Pods in a Deployment come and go, and their IP addresses change. So, it doesn't make sense to use Pod IP addresses directly. With a Service, you get a stable IP address that lasts for the life of the Service, even as the IP addresses of the member Pods change.

A Service also provides load balancing. Clients call a single, stable IP address, and their requests are balanced across the Pods that are members of the Service.

- How Pods are Connected to Service?

A Service identifies its member Pods with a selector. For a Pod to be a member of the Service, the Pod must have all of the labels specified in the selector. A label is an arbitrary key/value pair that is attached to an object.

Types of Services

There are five types of Services:

- 1) ClusterIP (default): Internal clients send requests to a stable internal IP address.
- 2) NodePort: Clients send requests to the IP address of a node on one or more nodePort values that are specified by the Service.
- 3) LoadBalancer: Clients send requests to the IP address of a network load balancer.
- 4) ExternalName: Internal clients use the DNS name of a Service as an alias for an external DNS name.
- 5) Headless: You can use a headless service in situations where you want a Pod grouping, but don't need a stable IP address.

The NodePort type is an extension of the ClusterIP type. So, a Service of type NodePort has a cluster IP address.

The LoadBalancer type is an extension of the NodePort type. So, a Service of type LoadBalancer has a cluster IP address and one or more nodePort values

##ClusterIP (default): Internal clients send request to a stable internal IP address.##

This will be used for internal pod communication only, kubernetes private ip communicate with cluster service, so it will use for internal communication, to communicate internal resources and not to external.

kubectl get nodes -o wide

```
[root@anskube manifest]# kubectl get nodes -o wide
```

NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	EXTERNAL-IP	OS-IMAGE	KE
gke-robo-default-pool-df26d11d-jl0j	Ready	<none>	11h	v1.13.11-gke.9	10.128.0.33	35.192.148.161	Container-Optimized OS from Google	4.
gke-robo-default-pool-df26d11d-kbk2	Ready	<none>	11h	v1.13.11-gke.9	10.128.0.34	35.232.14.42	Container-Optimized OS from Google	4.
gke-robo-default-pool-df26d11d-svhh	Ready	<none>	11h	v1.13.11-gke.9	10.128.0.35	35.222.31.76	Container-Optimized OS from Google	4.

@Create a pod.

#cat pods.yml

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-web1
  labels:
    app: web
spec:
  containers:
  - name: website1
    image: nginx:1.16
    ports:
    - containerPort: 80
```

kubectl apply -f pods.yml

kubectl get pods -o wide

```
[root@anskube manifest]# kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
nginx-web1	1/1	Running	0	2m19s	10.32.1.9	gke-robo-default-pool-df26d11d-svhh	<none>	<none>

#kubectl exec -it nginx-web1 -- nginx -v

```
[root@anskube manifest]# kubectl exec -it nginx-web1 -- nginx -v
nginx version: nginx/1.16.1
```

@@To check nginx http access between the containers, we are going to run ubuntu image.

kubectl run -it server1 --image=ubuntu /bin/bash

:- after login to pod, run below commands.

apt-get update

apt-get install curl

curl <http://10.32.1.9>

```
root@server1-578d4f748-rqshv:/# curl http://10.32.1.9
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
```

....

@@ Create cluster IP service.

#cat clusterport.yml

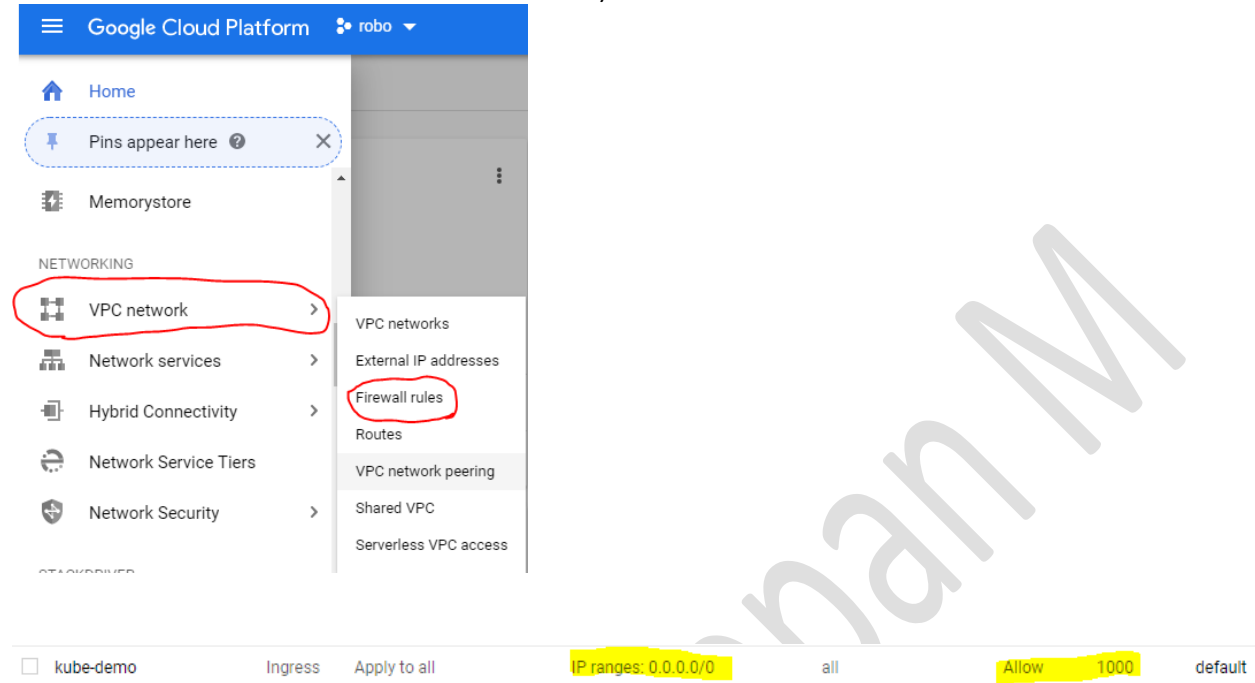
```
apiVersion: v1
kind: Service
metadata:
  name: service-clusterip1
spec:
  selector:
    app: web
  type: ClusterIP
  ports:
  - protocol: TCP
    port: 80
    targetPort: 80
```

```
# kubectl get svc
```

```
# kubectl describe svc service-clusterip1
```

In Kubernetes cluster, on each node, specific port will be opened with that port, we can expose the running application to external client access. on specific node, doesn't matter whether pod is running or not, however the service will run because of kubeproxy. the disadvantage is we access the node via specific port via public ip of VM's and we cannot ask end user to access with port and possibilities are there for hacking via port.

Note:- Make sure to create a firewall rules on GCP (Google Cloud Platform) for node port access (VPC network --> Firewall rules --> create firewall rules)



Type1:- Creating Nodeport Service with Dynamic port.

cat nodeport.yml

```
apiVersion: v1
kind: Service
metadata:
  name: my-np-service
spec:
  selector:
    app: web
  type: NodePort
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

kubectl apply -f nodeport.yml

```
[root@anskube manifest]# kubectl get svc
NAME                TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
kubernetes           ClusterIP   10.36.0.1     <none>          443/TCP          3h58m
my-np-service        NodePort    10.36.7.82    <none>          80:30552/TCP     92s
service-clusterip1   ClusterIP   10.36.3.26    <none>          80/TCP           150m
```

@Try to access the website by using external ip with auto generate port.

kubectl get pods --show-labels

kubectl get nodes -o wide

:-Open browser and access with dynamic port <http://35.192.148.161:30552/>

Type2:- Creating Nodeport Service with static port.

```
# cat nodeport1.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: my-nodeport-service
spec:
  selector:
    app: web
  type: NodePort
  ports:
  - name: http
    protocol: TCP
    port: 80
    targetPort: 80
    nodePort: 30036
```

```
# kubectl apply -f nodeport1.yml
```

```
# kubectl get svc
```

```
[root@anskube manifest]# kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.36.0.1	<none>	443/TCP	4h12m
my-nodeport-service	NodePort	10.36.12.21	<none>	80:30036/TCP	69s

@Try to access the website by using external ip with specific port.

<http://35.192.148.161:30036/>

##LoadBalancer: Clients send requests to the IP address of a network load balancer##

This will come up with single static ip and listen via 80 port and backend whatever pod running, it will pass the request from the client. so, end user will access the domain without port mention, and this will be used in real-time. the disadvantage is incased four application has exposed in the services means, four load balancers will be created and will be more cost involved, so these kind of setup will be used for cloud providers.

Type1:- In this method, gcloud provider will assign the external ip to load balancer, we can access via that ip.

```
# cat loadbalancer.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: my-nlb-services
spec:
  selector:
    app: web
  type: LoadBalancer
  ports:
  - port: 80
    targetPort: 80
```

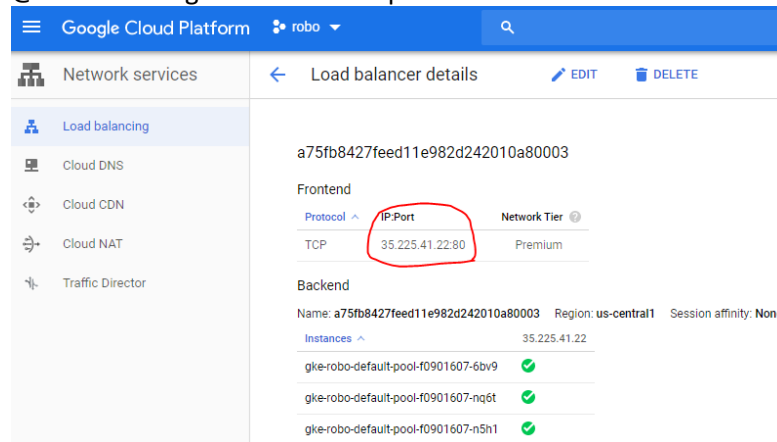
```
# kubectl apply -f loadbalancer.yml
```

```
# kubectl get svc -o wide
```

```
[root@anskube manifest]# kubectl get svc -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
kubernetes	ClusterIP	10.36.0.1	<none>	443/TCP	4h28m	<none>
my-nlb-services	LoadBalancer	10.36.15.171	35.225.41.22	80:30976/TCP	3m53s	app=web

@You can see get the external ip info from GCP as well.



Google Cloud Platform

Network services

Load balancing

Cloud DNS

Cloud CDN

Cloud NAT

Traffic Director

Load balancer details

EDIT DELETE

a75fb8427feed11e982d242010a80003

Frontend

Protocol	IP:Port	Network Tier
TCP	35.225.41.22:80	Premium

Backend

Name: a75fb8427feed11e982d242010a80003 Region: us-central1 Session affinity: None

Instances

Instance	Status
gke-rob-default-pool-f0901607-6bv9	✓
gke-rob-default-pool-f0901607-nq6t	✓
gke-rob-default-pool-f0901607-n5h1	✓

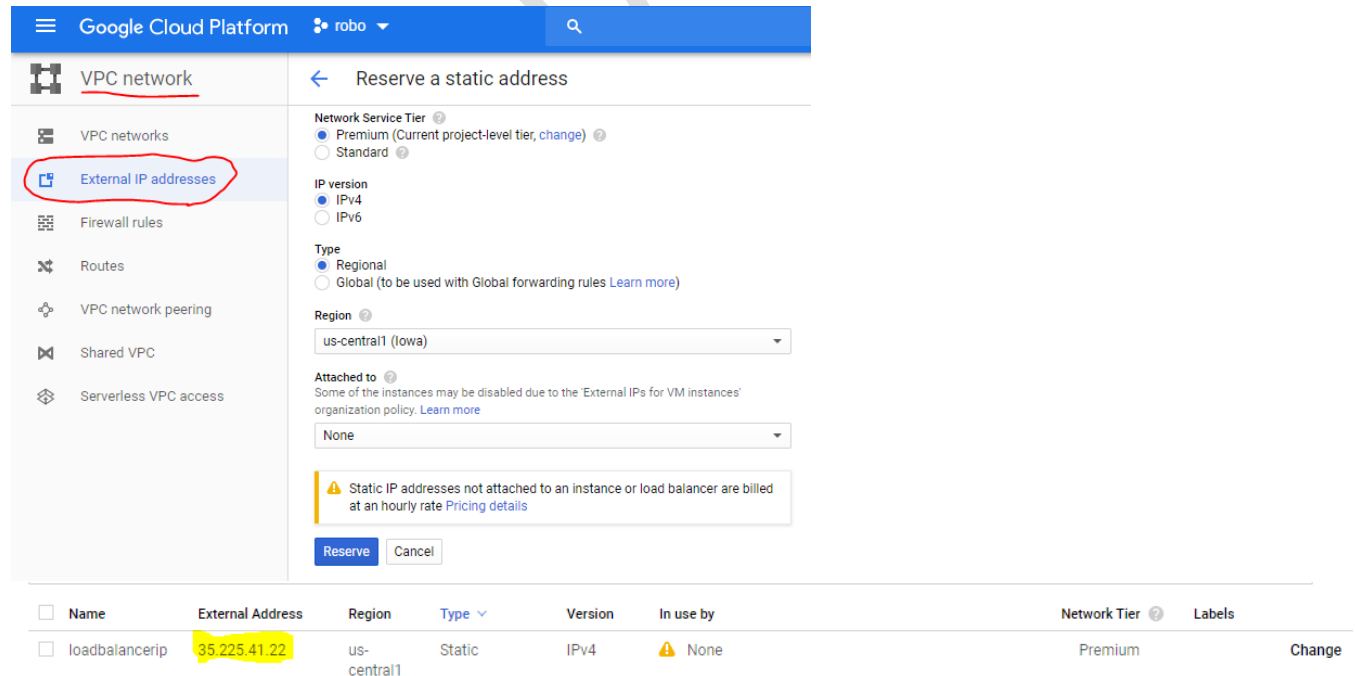
:-Open browser and try to access with loadbalancer ip. <http://35.225.41.22/>

Note:- Delete once you done.

kubectl delete -f loadbalancer.yml

Type2:- In this method, we can use static ip which is preferable for prod, since dynamic ip can be changed by cloud provider and it may disconnect the access.

Note:- Reserve the static ip on GCP cloud by selecting VPC networks --> External IP addresses --> RESERVE STATIC ADDRESS



Google Cloud Platform

VPC network

VPC networks

External IP addresses

Firewall rules

Routes

VPC network peering

Shared VPC

Serverless VPC access

Reserve a static address

Network Service Tier

☒ Premium (Current project-level tier, change)

☐ Standard

IP version

☒ IPv4

☐ IPv6

Type

☒ Regional

☐ Global (to be used with Global forwarding rules Learn more)

Region

us-central1 (Iowa)

Attached to

None

Static IP addresses not attached to an instance or load balancer are billed at an hourly rate Pricing details

Reserve Cancel

Name	External Address	Region	Type	Version	In use by	Network Tier	Labels
loadbalancerip	35.225.41.22	us-central1	Static	IPv4	None	Premium	Change

:- with this reserve external ip, we are going to create load balancer on Kubernetes cluster.

```
#cat loadbalancer1.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: my-nlb-services
spec:
  selector:
    app: web
  type: LoadBalancer
  loadBalancerIP: 35.225.41.22
  ports:
    - port: 80
      targetPort: 80
```

```
# kubectl apply -f loadbalancer1.yml
```

```
# kubectl get svc
```

```
[root@anskube manifest]# kubectl get svc
NAME                TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
kubernetes          ClusterIP     10.36.0.1     <none>         443/TCP          4h45m
my-nlb-services     LoadBalancer 10.36.2.173   35.225.41.22  80:30046/TCP    110s
```

:- Open browser and try to access via LB ip. <http://35.225.41.22/>

##ExternalName: Internal clients use the DNS name of a Service as an alias for an external DNS name##
for example, our own machine will connect to third-party application is called external service and which will be mapped in this external service and it has an instruct to access . every time no need to do any changes on application level and since external ip has given by apps vendor, so it may change at any time and not required to change on all backend pods. Here we can make changes on external name which we configured and that will manage backend.

```
# cat externlsvc.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: ExternalName
  externalName: yahoo.com
```

```
# kubectl apply -f externlsvc.yml
```

```
# kubectl get svc
```

```
[root@anskube manifest]# kubectl get svc
NAME                TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
kubernetes          ClusterIP     10.36.0.1     <none>         443/TCP          4h58m
my-nlb-services     LoadBalancer 10.36.2.173   35.225.41.22  80:30046/TCP    15m
my-nodeport-service NodePort      10.36.12.21   <none>         80:30036/TCP    47m
my-np-service       NodePort      10.36.7.82    <none>         80:30552/TCP    61m
my-service          ExternalName   <none>        yahoo.com      <none>          14s
```

```
# kubectl get pods
```

@Login to pod and ping with external name.

```
# kubectl exec -it nginx-web1 /bin/bash
```

```
#apt-get update && apt-get install iputils-ping → Run these commands inside the pod.
```

```
root@nginx-web1:/# ping my-service
PING yahoo.com (72.30.35.10) 56(84) bytes of data.
64 bytes from media-router-fp2.prod1.media.vip.bf1.yahoo.com (72.30.35.10): icmp_seq=1 ttl=46 time=43.10 ms
64 bytes from media-router-fp2.prod1.media.vip.bf1.yahoo.com (72.30.35.10): icmp_seq=2 ttl=46 time=43.10 ms
```